AI Assistant

Take a pic of a Question Paper (previous years obv) and upload here:

Choose an image

Drag and drop file here

Limit 200MB per file • JPG, JPEG, PNG

Browse files

Scroll down to Download the AI Generated QnA docx file.

PDF version coming soon In Sha' Allah 😉

The questions detected are:

ne, — wang (con:

. a) What is the trade off between bias and variance? b) Differentiate classification and regression tasks with examples. c) What is imbalanced dataset? How would you handle it? d) Why XOR problem cannot be solved by a single layer perceptron?® » e) Define Markov decision process. xf) Discuss the importance of pooling layer in convolutional neural networks. g) What is Fi score? Explain its significance.

Generating answers...

The Answers generated are:

1. What is the trade off between bias and variance?

Bias is the error made by a model when it consistently underestimates or overestimates the true value. Variance is the error made by a model when it is sensitive to small changes in the training data. A model with high bias will be inaccurate, but a model with high variance will be unstable. The goal of machine learning is to find a model that minimizes both bias and variance.

2. Differentiate classification and regression tasks with examples.

Classification is a task in which the model is asked to predict a class label for each data point. For example, a model might be asked to predict whether a picture contains a cat or a dog. Regression is a task in which the model is asked to predict a continuous value for each data point. For example, a model might be asked to predict the price of a house.

3. What is imbalanced dataset? How would you handle it?

An imbalanced dataset is a dataset in which there are more data points in one class than in another. This can be a problem for machine learning models, because they can be biased towards the majority class. There are a number of ways to handle imbalanced datasets, such as oversampling the minority class, undersampling the majority class, or using a cost-sensitive learning algorithm.

4. Why XOR problem cannot be solved by a single layer perceptron?

A single layer perceptron is a neural network with one hidden layer. It can be used to solve classification problems, but it cannot be used to solve regression problems. The XOR problem is a Boolean function that takes two inputs and outputs 1 if the inputs are different and 0 if the inputs are the same. It cannot be solved by a single layer perceptron because the perceptron cannot learn a non-linear function.

5. Define Markov decision process.

A Markov decision process (MDP) is a mathematical model that describes a decision-making process under uncertainty. It is a state-action-reward model, in which the state of the system at time t is denoted by s_t, the action taken at time t is denoted by a_t, and the reward received at time t is denoted by r_t. The state of the system at time $t+1$ is determined by the state of the system at time t and the action taken at time t, according to the transition probability function $p(s_{t+1}|s_t, a_t)$. The reward received at time t is determined by the state of the system at time t and the action taken at time t, according to the reward function $r(s_t, a_t)$. The goal of an MDP is to find a policy π that maximizes the expected reward over time.

6. Discuss the importance of pooling layer in convolutional neural networks.

A pooling layer is a layer in a convolutional neural network that downsamples the input data. This is done by taking a small neighborhood of the input data and averaging the values in the neighborhood. This has the effect of reducing the size of the input data, which can help to improve the performance of the network. Pooling layers are also used to reduce the number of parameters in the network, which can help to improve the speed of training.

7. What is Fi score? Explain its significance.

The F1 score is a measure of the performance of a binary classifier. It is calculated as the harmonic mean of the precision and recall. The precision is the fraction of positive predictions that are correct, and the recall is the fraction of positive examples that are correctly predicted. The F1 score is a good measure of the overall performance of a classifier, because it takes into account both precision and recall.

Download Docx with Answers